

Logging the right way!

Devi A S L

PyCon India, 2020



About

- @asldevi
- 15 years of experience in industry
- Staff Engineer at Razorpay
- Ex-Architect @ PowerToFly





Why Logging?

We want to find out what had happened

Why Logging?

We want to find out what **exactly** had happened

Why Logging?

We want to find out what **exactly** had happened **quickly**

Why Logging?

We want to find out what **exactly** had happened **quickly** and it's **cause**



Agenda

- Making logs easily searchable, filterable and aggregate-able
- Balancing debuggability and logging cost
- Tracing requests across microservices in logs

Sample Log Line



```
logger.info("Transaction succeeded with id %s, amount  
%f, payment mode %s", txn["id"], txn["amount"],  
txn["mode"])
```


Sample Log Line

```
logger.info("Transaction succeeded with id %s, amount  
%f, payment mode %s", txn["id"], txn["amount"],  
txn["mode"])
```

Searchability

[app][info] 2020-10-02T18:23:13.972808: Transaction succeeded with id Fbgi4yFzwLc4NQ, amount 1025.00, payment mode NEFT

- Human readable, not machine readable
- Filter searches *
- Aggregates *

* Unless your log management system is doing the heavy lifting

Treat logs as event streams



Enter Structured Logging

```
logger.info("Transaction succeeded with id %s, amount  
%f, payment mode %s", txn["id"], txn["amount"],  
txn["mode"])
```

```
logger.info("new_transaction", id=txn['id'],  
amount=txn['amount'], payment_mode=txn['mode'],  
is_successful=True)
```

Structured Logging

```
{  
  message: 'new_transaction',  
  id: 'Fbgi4yFzwLc4NQ',  
  amount: 1025.00,  
  payment_mode: 'NEFT',  
  is_successful: True,  
  timestamp: 2020-10-02T18:23:13.972808  
}
```

- Filter searches ✓
- Aggregates ✓

Structured Logging with Structlog

- Drop in replacement to standard library
 - `s/logging.getLogger/structlog.get_logger`
- Plenty of powerful pipelines
 - add metadata, redact sensitive info etc

Logging is not Cheap

Number of logs \propto number of requests

Cost \propto number of logs * cost of (network + storage) * retention period

Logging is not Cheap

At Razorpay, we have 2 TB logs/day flowing in



Optimize for cost or debuggability?

Dynamic logging



Change Log Level in Run time

```
logging.config.listen(<port>)
```

Works with fileConfig, dictConfig etc

Change log level in run time

```
logger = structlog.getLogger(...)
```

```
logger.setLevel(level)
```




can be changed
run-time
based on metrics
errors/latencies

Logging microservices



Request flow and logs



```
{
  message: log_line_from_frontend,
  timestamp: 91
}
{
  message: log_line_from_frontend,
  timestamp: 92
}
{
  message: log_line_from_A,
  timestamp: 93
}
{
  message: log_line_from_A,
  Timestamp: 94
}
{
  message: log_line_from_B,
  timestamp: 95
}
```

Tracing requests end-to-end

```
{
  message: log_line_from_frontend,
  timestamp: 91
}
{
  message: log_line_from_frontend,
  timestamp: 92
}
{
  message: log_line_from_A,
  timestamp: 93
}
{
  message: log_line_from_A,
  Timestamp: 94
}
{
  message: log_line_from_B,
  timestamp: 95
}
```



```
{
  message: log_line_from_frontend,
  timestamp: 91, request_id: abc
}
{
  message: log_line_from_frontend,
  timestamp: 92, request_id: xyz
}
{
  message: log_line_from_D,
  timestamp: 93, request_id: abc
}
{
  message: log_line_from_D,
  Timestamp: 94, request_id: def
}
{
  message: log_line_from_A,
  timestamp: 95, request_id: abc
}
```


Tracing requests end-to-end

```
{
  message: log_line_from_frontend, ...
  timestamp: 91
}
{
  message: log_line_from_frontend, ...
  timestamp: 92
}
{
  message: log_line_from_A, ...
  timestamp: 93
}
{
  message: log_line_from_A, ...
  Timestamp: 94
}
{
  message: log_line_from_B, ...
  timestamp: 95
}
```



```
{
  message: log_line_from_frontend, ...
  timestamp: 91, request_id: abc
}
{
  message: log_line_from_frontend, ...
  timestamp: 92, request_id: xyz
}
{
  message: log_line_from_A, ...
  timestamp: 93, request_id: abc
}
{
  message: log_line_from_A, ...
  Timestamp: 94, request_id: def
}
{
  message: log_line_from_B, ...
  timestamp: 95, request_id: abc
}
```

Request ID through Header propagation



Recap

- Set log levels appropriately, put your own if required
- Machine readable logs aka structured logs - always
- Dynamic logging to have control of log volume
- Unique correlation id across all microservices connecting them all

References

- <https://docs.python.org/2/howto/logging-cookbook.html#configuration-server-example> (is similar in py3)
- <https://www.structlog.org/en/stable/index.html>
- <https://medium.com/hiredscore-engineering/logging-lets-do-it-right-41d568d3bfcd>



Questions?

Grow with us

@asldevi

